



NSDb

Leveraging Scala and
Akka to build NSDb,
a distributed time-series database

Saverio Veltri @save_veltri Firenze
Paolo Mascetti @mascettipaolo 14th September

Who we are



Saverio Veltri
Solution Architect



Paolo Mascetti
Data Engineer



RADICALBIT

- Based in Milan since 2015
- Event Stream Processing products and solutions

We are a specialized software firm,
born in Milan on 2015





RADICALBIT

- Based in Milan since 2015
- Event Stream Processing products and solutions

We are focussed on the design and development of **Event Stream Processing** products and solutions, combining **streaming technologies** with **Machine Learning** and **A.I.**



Agenda

Introduction

NSDb Main Features

Single Node Design

Akka Cluster Overview

Distributed Design

Roadmap & Licensing

Contribution



Introduction

Motivations

Connotations

Time Series Model

Consistency Model

NSDb in Data Intensive Architectures

NSDb in CQRS Pattern



Motivations

- Have a deep technical ownership of the solution
- Too many licensing and pricing issues exploring third-party OEM solutions
- Third-party solutions don't completely fit our requirements

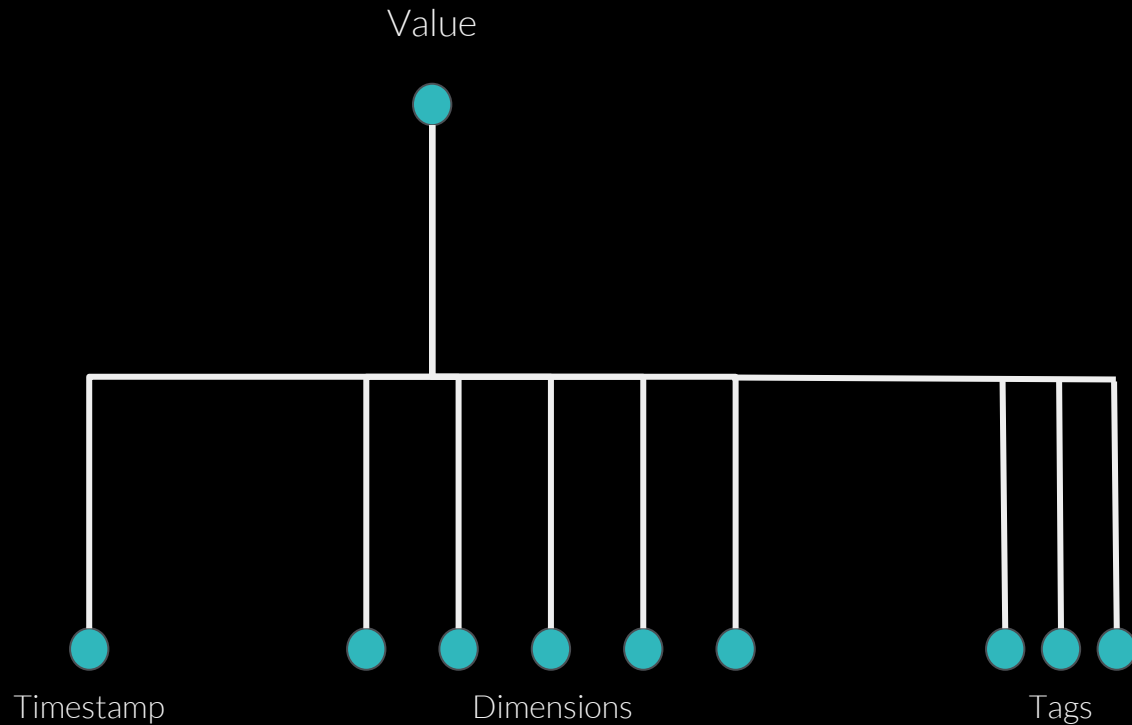
Connotations

- Distributed
 - Allows cluster deploy of p2p nodes
 - Based on Akka Cluster
- TimeSeries
 - Optimized time series management
- Streaming oriented
 - Maintain real-time capability in streaming architectures



Time Series Model (I)

Bit: a MultiDimensional Time Series value



Timestamp: the record time

Value: the numerical value being measured

Dimensions: a dynamic list of queryable String -> Value pairs

Tags: special dimensions user can apply aggregations on



Time Series Model (II)

- NSDB's Bits are **immutable**. New data continuously arrives, and will be always inserted and never updated.
- Bit schema is monotonic


Bit organization:

- **Metric**: a series of Bit (Records)
- **Namespace**: high level structure grouping metrics
- **Database**: logical container grouping namespaces



NSDb - Consistency Model

- Eventual consistency
- Real time delivery for subscribed client




NSDb in data intensive architectures

- Eventual Consistency narrows down the points of applicability of NSDb
- Real time streaming and Push features perfectly fit the serving layer (e.g. Kappa architecture and CQRS)



NSDb in CQRS Pattern

- Clear separation of Commands and Queries
- Scalability guaranteed by using 2 different databases



NSDb Main Features

NSDb Sharding

Natural Time Sharding

Data Partitioning

APIs & Connectors


Publish Subscribe



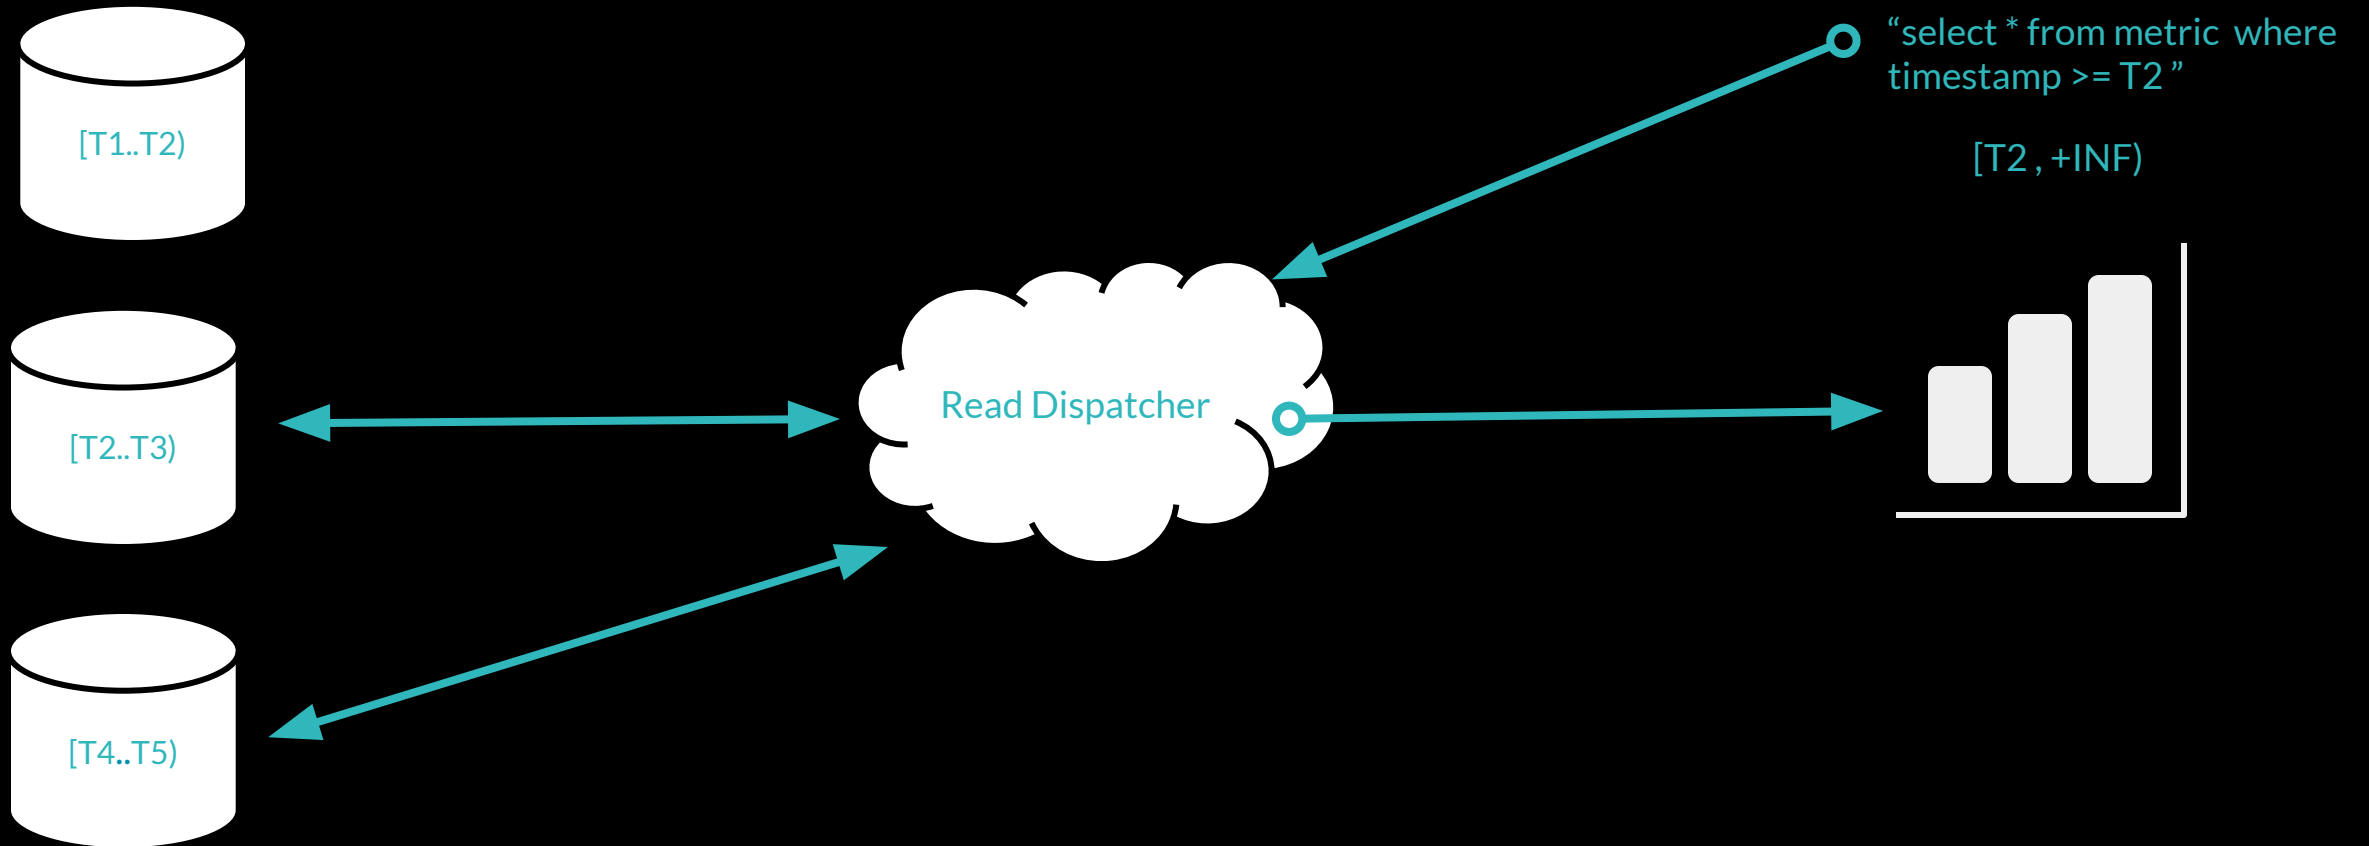
Natural Time Sharding

- Time Series points are gathered into Shards based on “event time”
- Any other partitioning will be demanded to Lucene indices
- This concept optimizes some time related frequent access patterns
- Data chunks are concatenated (and in case ordered) and not merged

Data Partitioning - Write



Data Partitioning - Read



APIs & Connectors

- Scala & Java APIs
- HTTP(S) APIs implemented using Akka HTTP
- WS APIs
- Flink Sink
- Kafka Connector



Scala Write APIs

```
implicit val executionContext: ExecutionContextExecutor = ExecutionContext.global
val NSDb =
  Await.result(NSDB.connect(host = "127.0.0.1", port = 7817), 10 seconds)
val series = NSDb
  .db(name = "conferences")
  .namespace(namespace = "Italy")
  .metric(metric = "scala-italy-attendees")
  .value(300)
  .dimension("city", "Florence")
  .tag("topic", "scala")
val res: Future[RPCInsertResult] = NSDb.write(series)
```



Scala Read APIs

```
implicit val executionContext: ExecutionContextExecutor = ExecutionContext.global
val connection = Await.result(NSDB.connect(host = "127.0.0.1", port = 7817), 10.seconds)

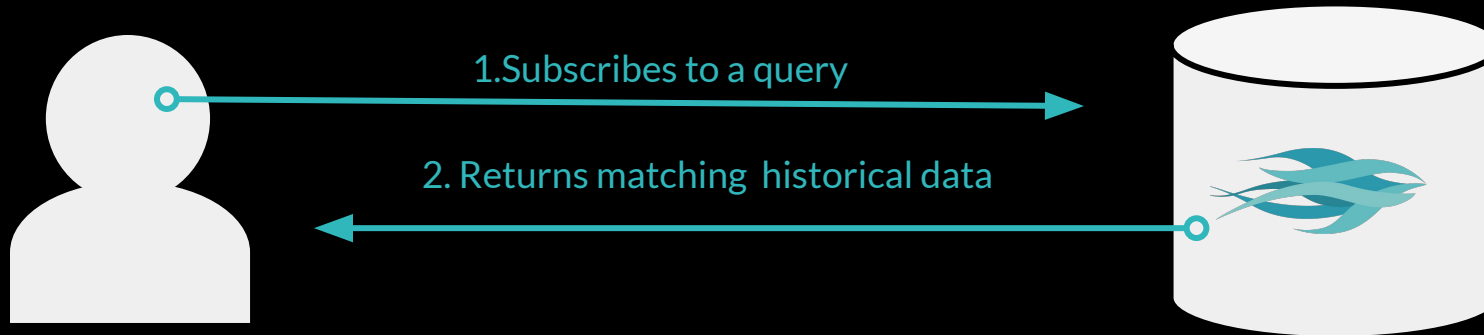
val query = connection
    .db( name = "conferences")
    .namespace( namespace = "italy")
    .query( queryString = "select * from scala-italy-attendees order by timestamp desc")

val readRes: Future[SQLStatementResponse] = connection.execute(query)
```



Publish-Subscribe (I)

1. User subscribes a query using WebSocket APIs
2. Historical data matching the query is returned



Publish-Subscribe (II)

3. Everytime new bits are written into NSDb, if they match user registered queries, are published on WebSocket channel



Single Node Design

Akka Recap

Overall Node Architecture

Lucene as Storage Layer

SQL Like Support

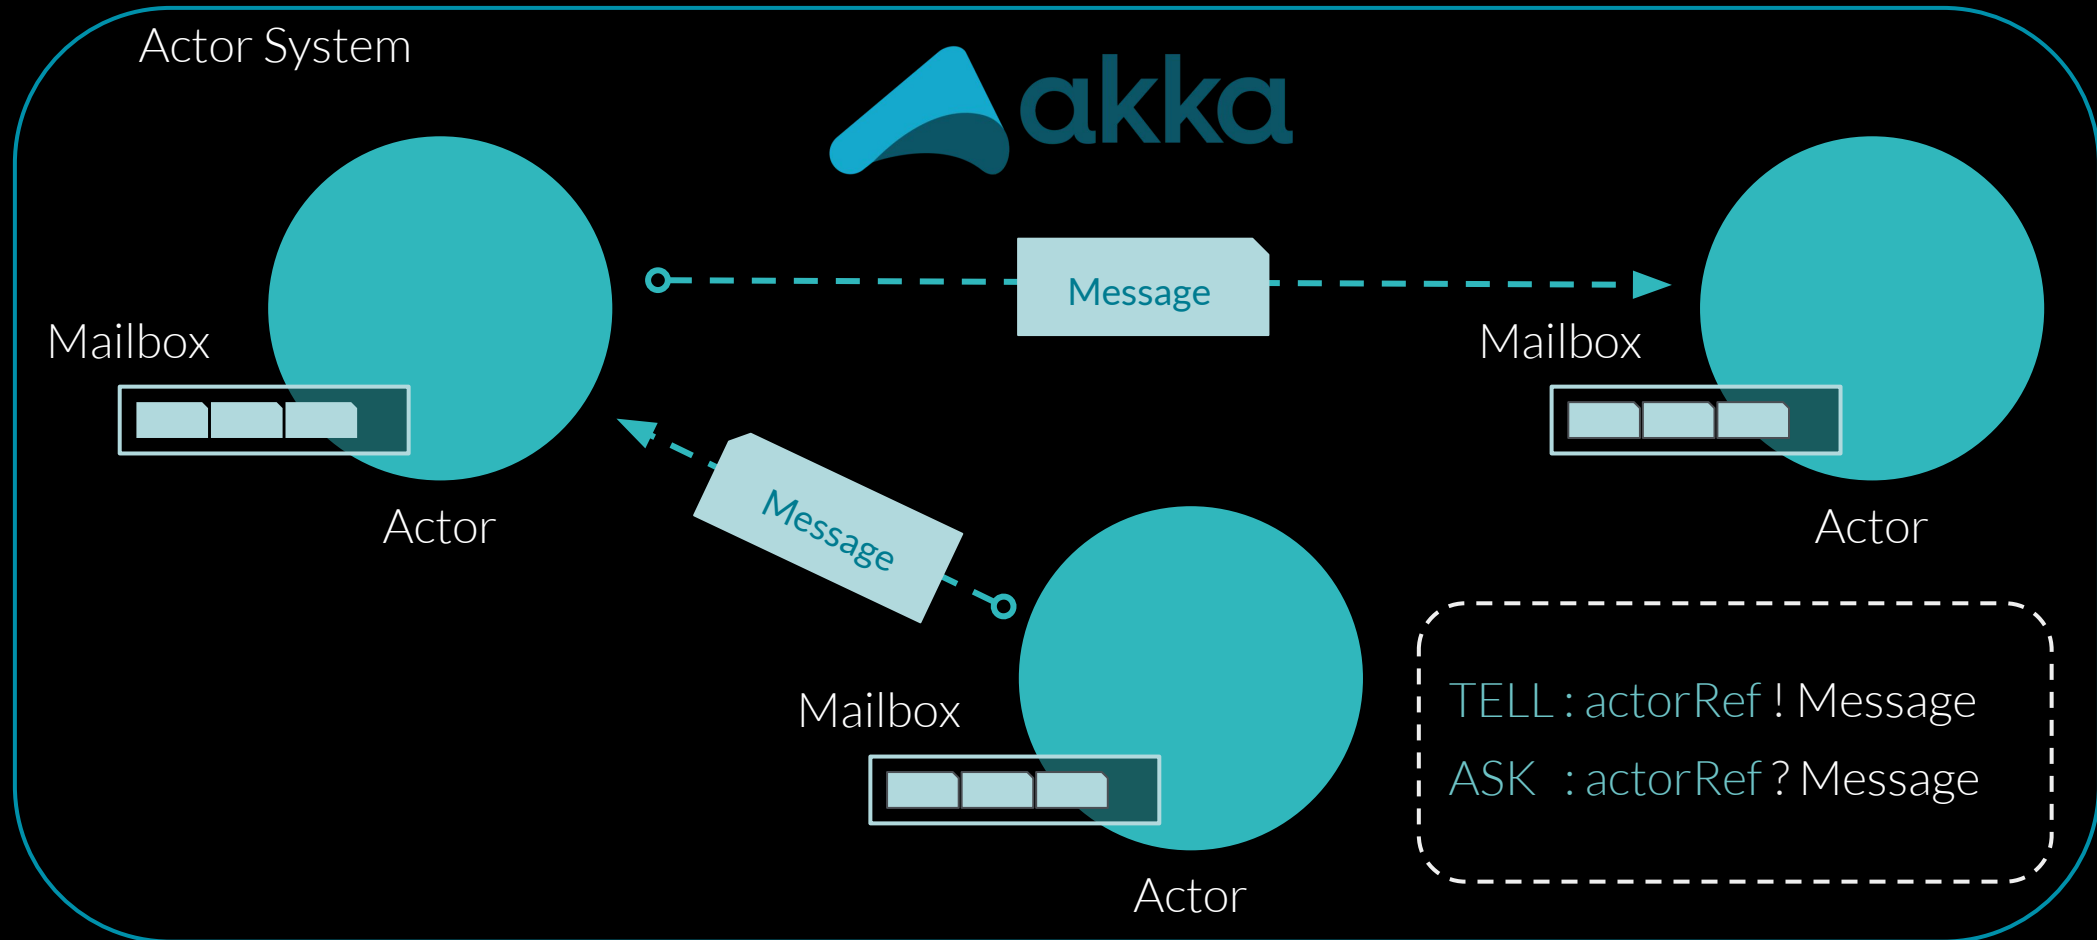
Handling mutable Lucene indices with
Akka

Node actors hierarchy

Data Streaming




Akka Recap (I)

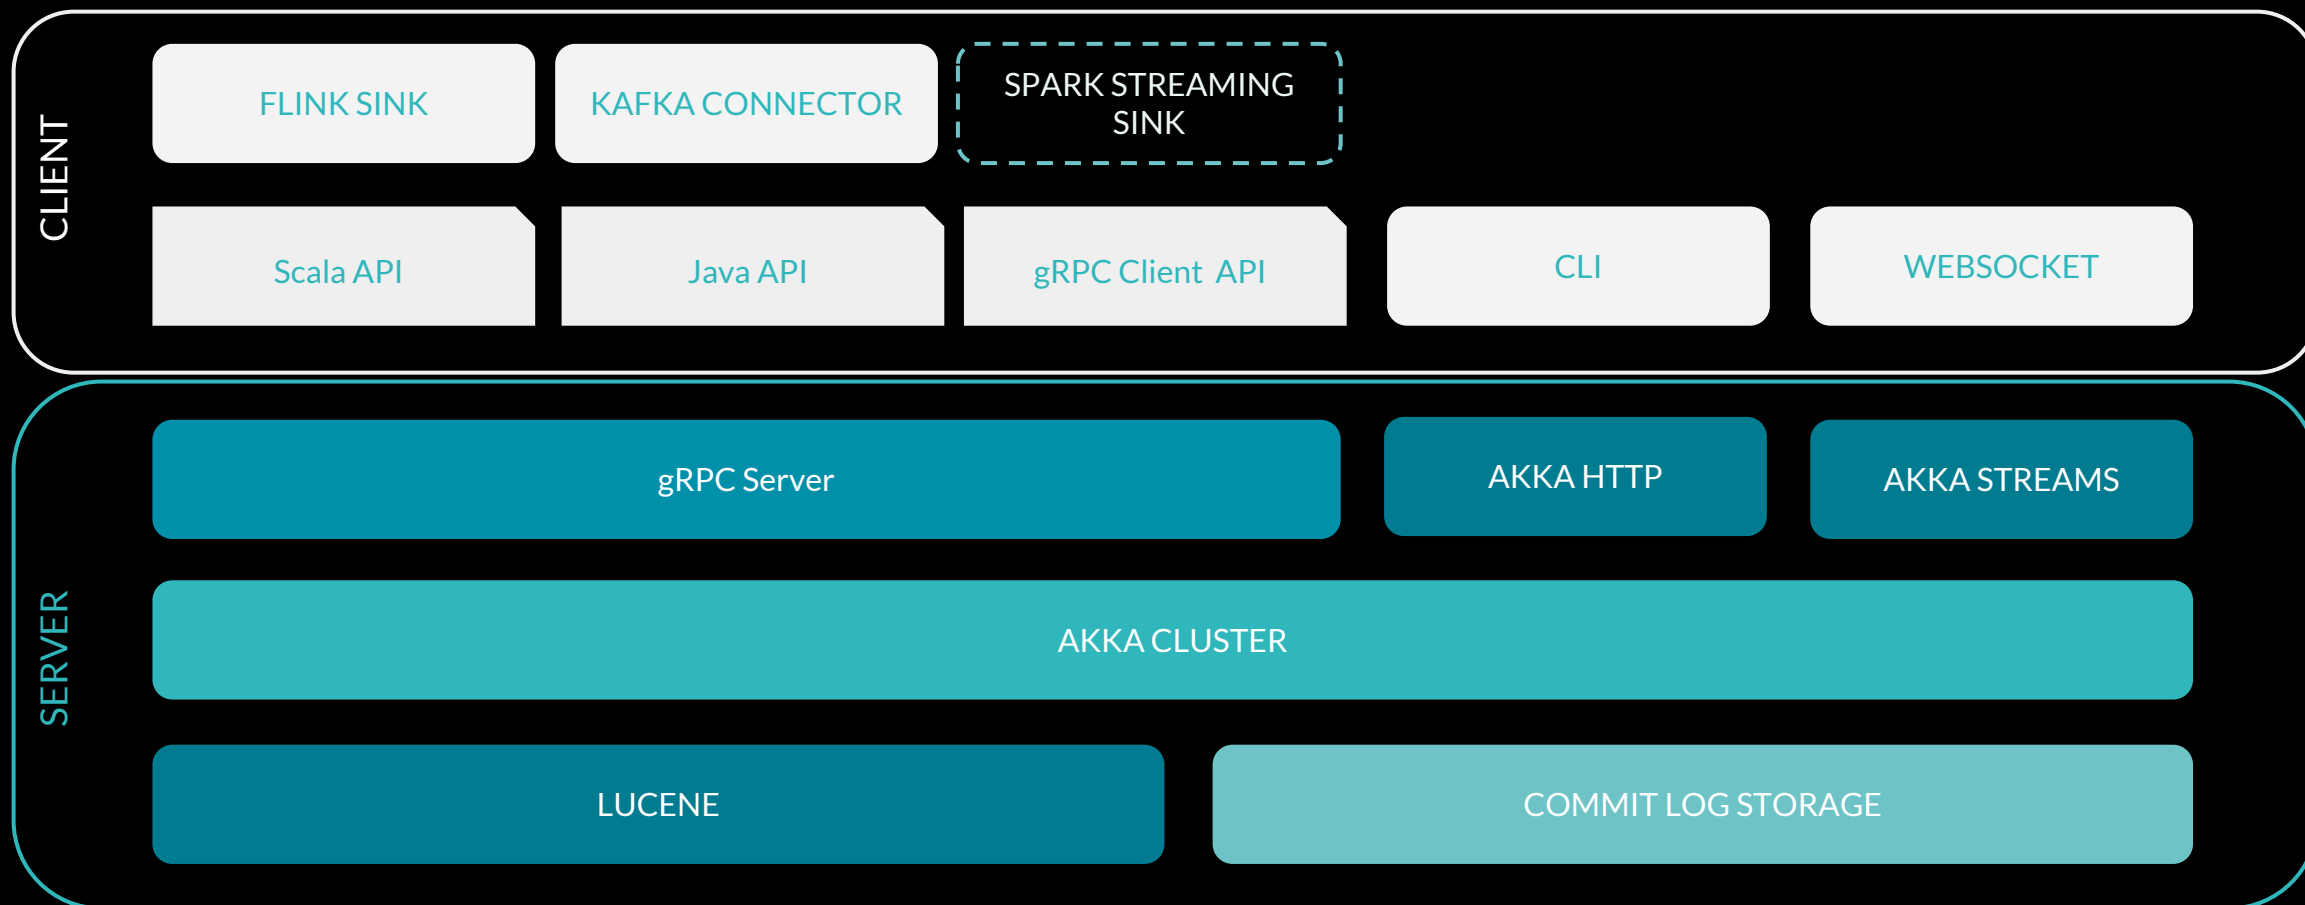


Akka Recap (II)

Actor System



Overall Node Architecture



Lucene as Storage Layer (I)



“Apache Lucene is an open source project implementing full-featured text search engine library written entirely in Java.”

- Ad Hoc indices management according to time-series handling



Lucene as Storage Layer (II)

PROs:

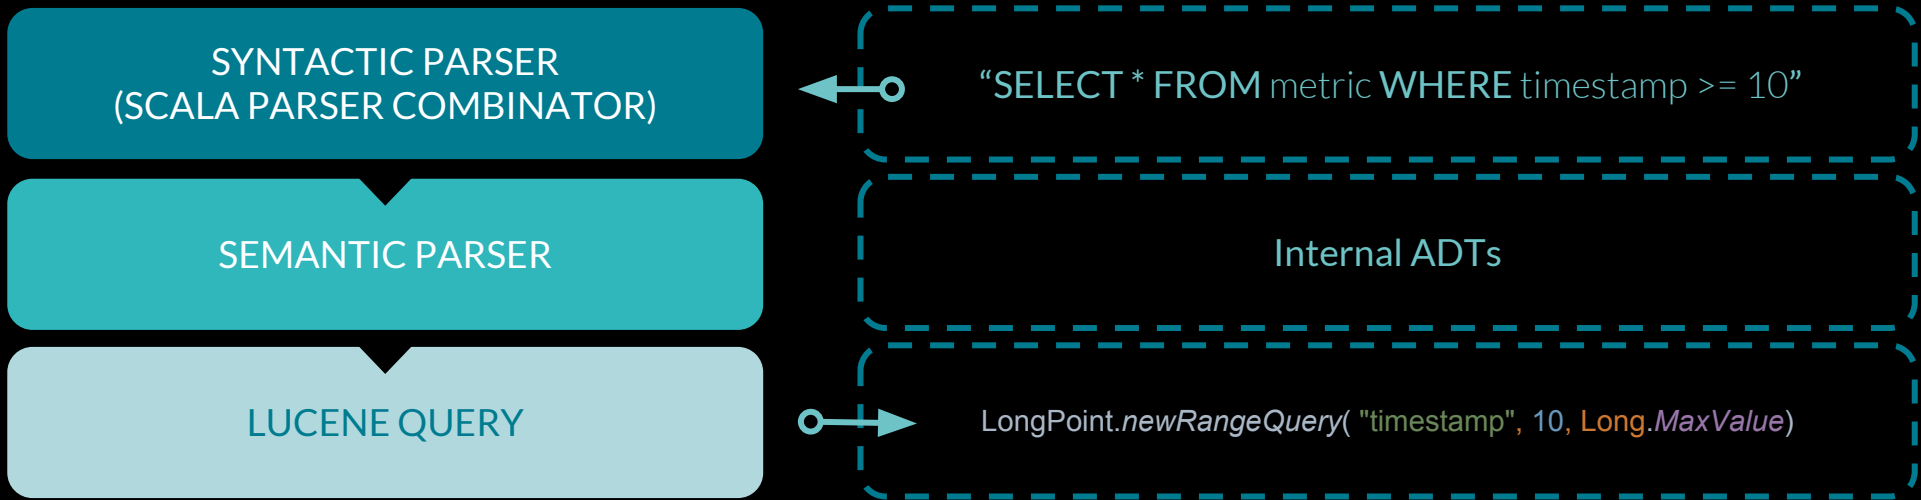
- Stable and continuously improved project
- Scalable, High-Performance Indexing
- Very common choice in database field
- Powerful query optimization
- Java implementation

CONs:

- Lack of documentation
- Java implementation



SQL Like Support

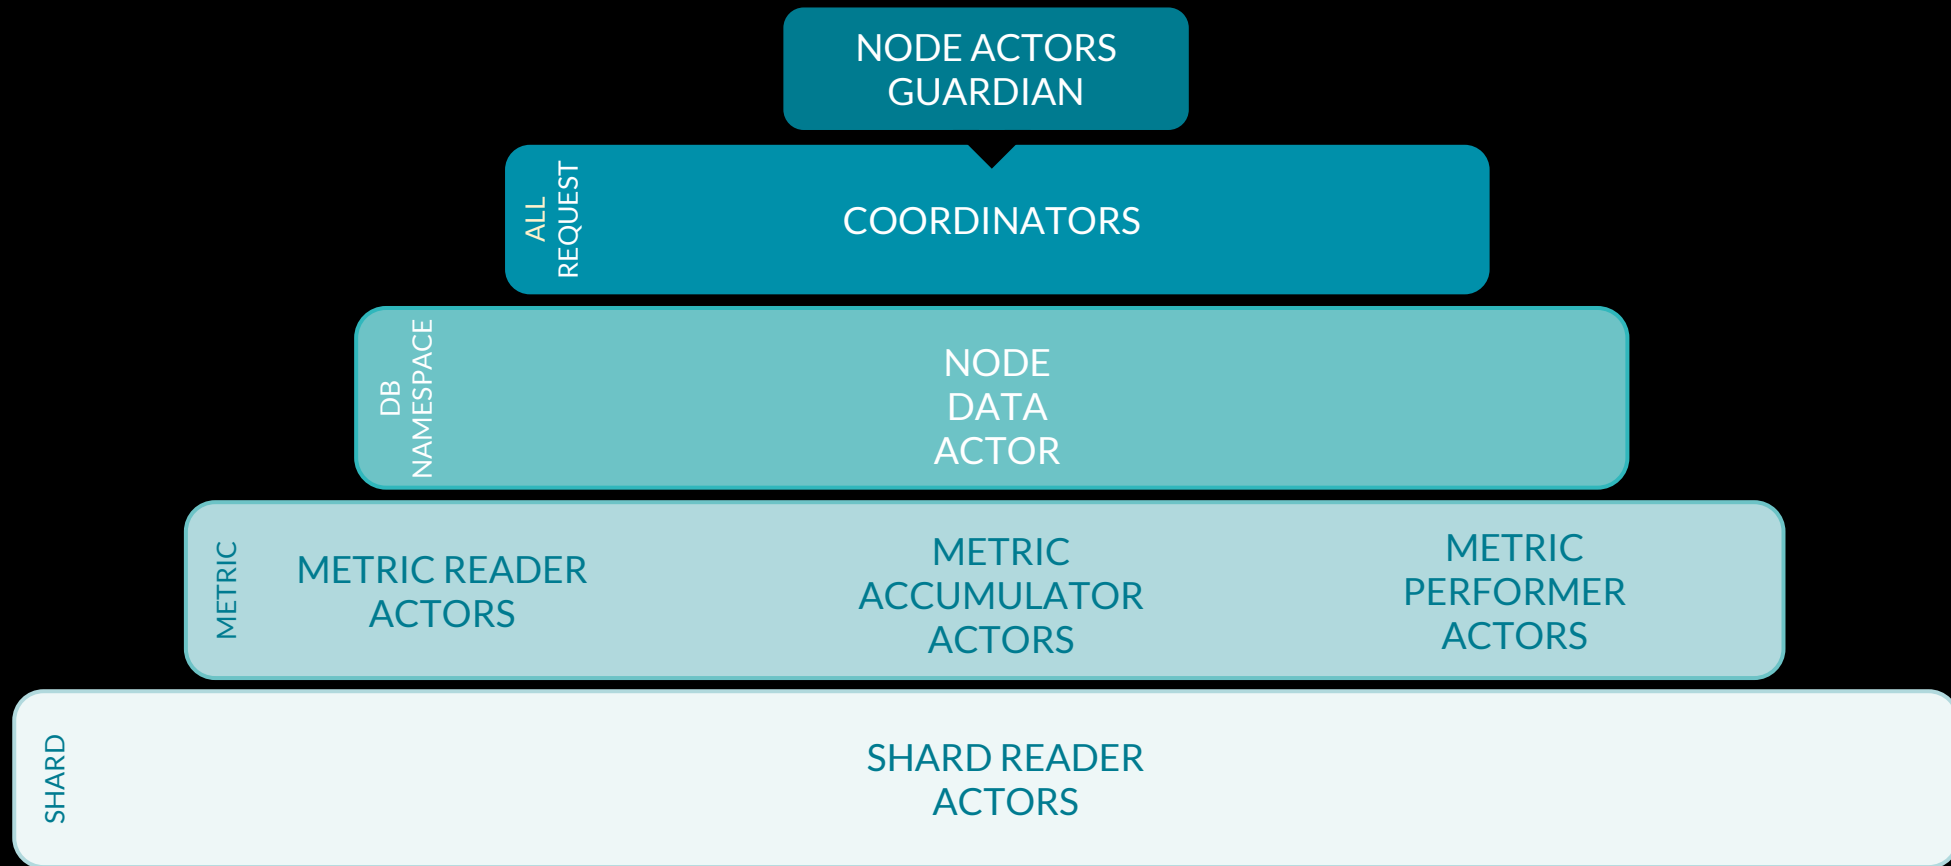


Handling mutable Lucene indices with Akka

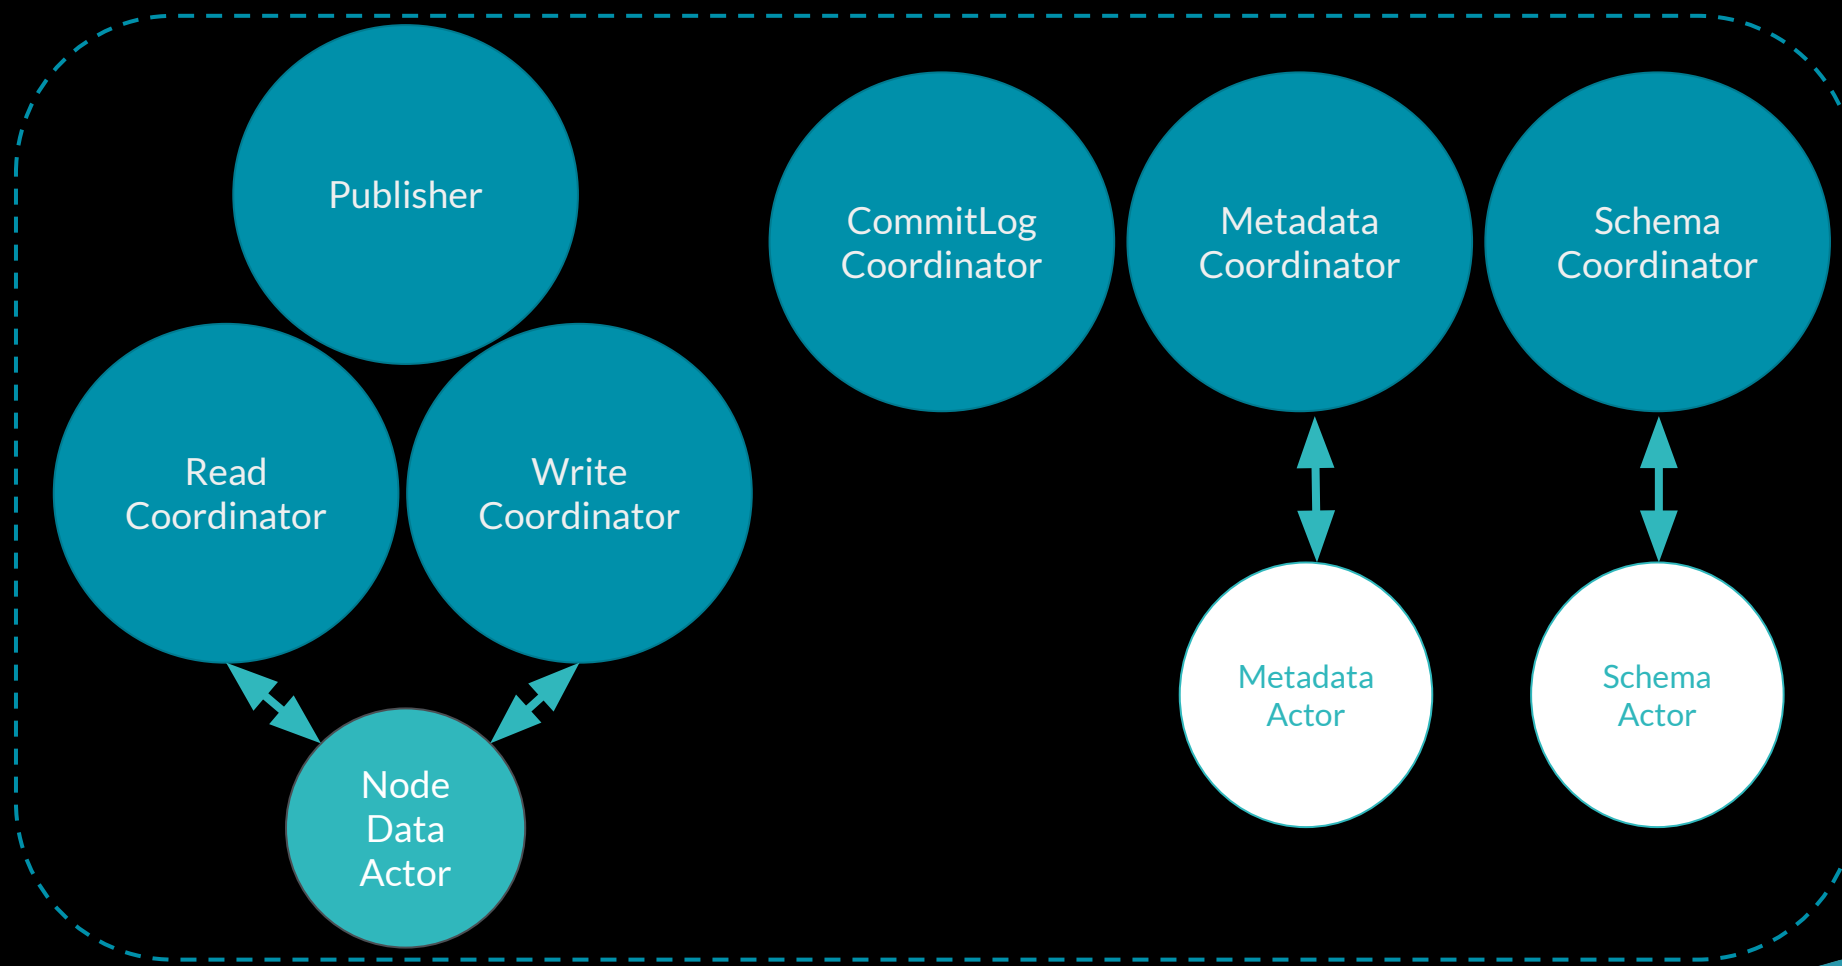
- Usage of message passing avoids locking and blocking
- Akka Actors wraps our own Lucene access layer
- Each Actor handles a single kind of operation (read or write) on a specific index
- Scale up on single node



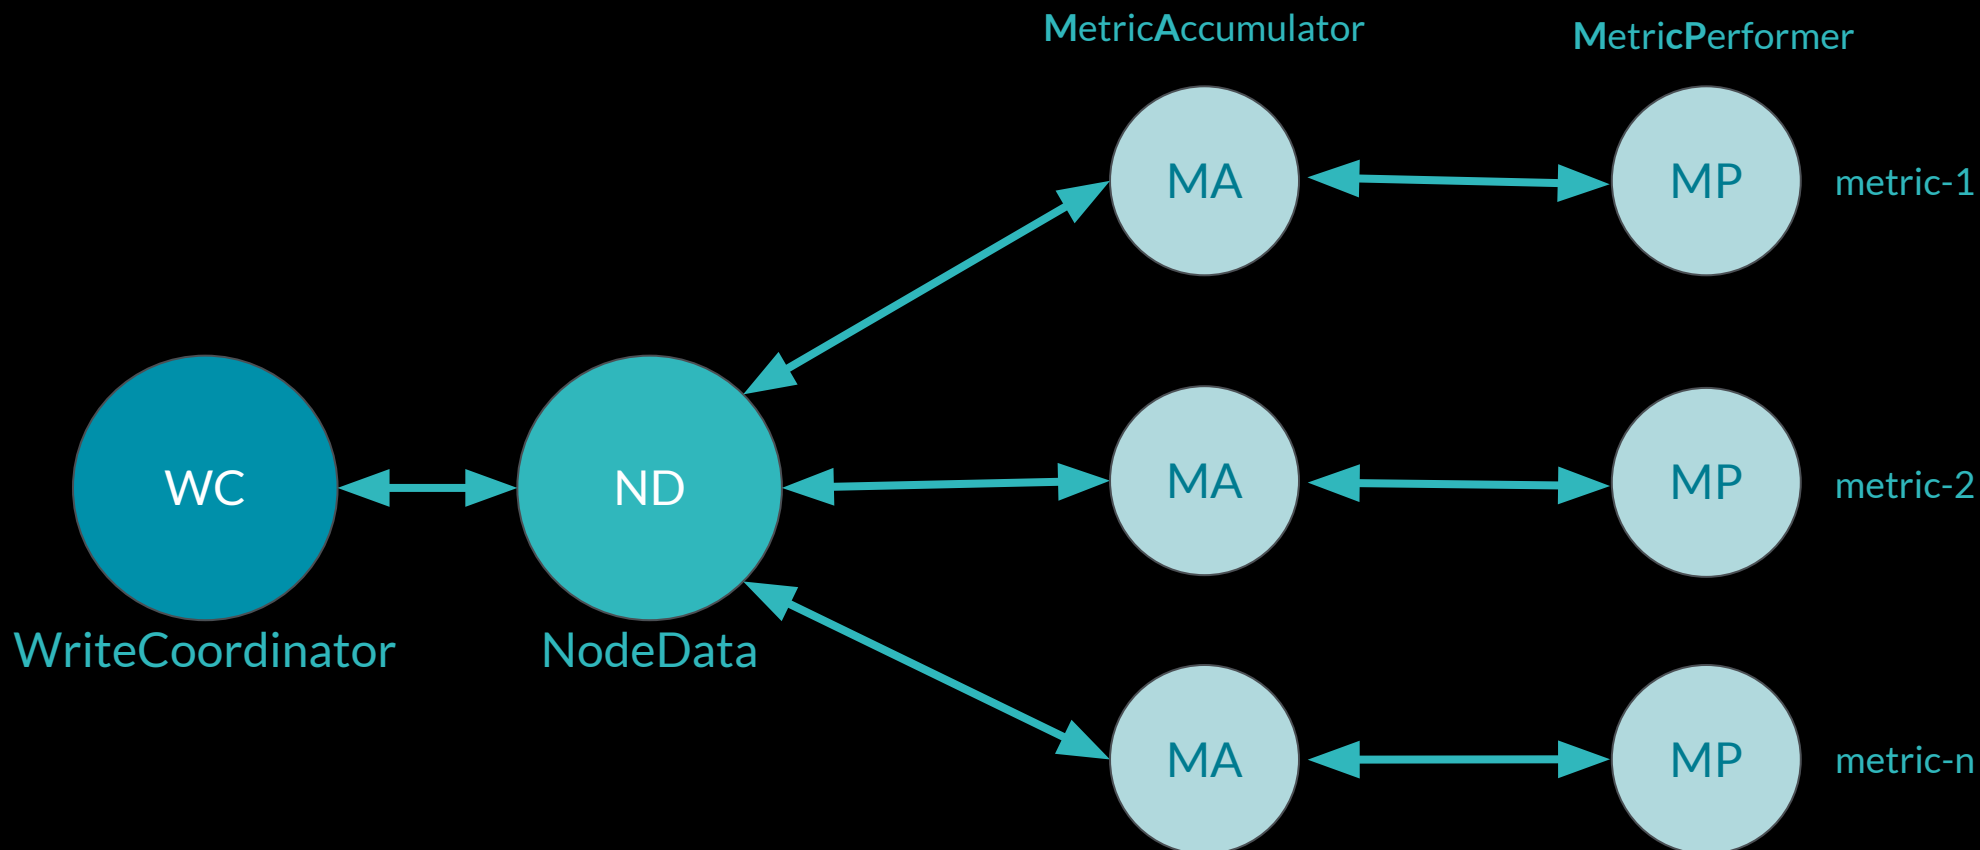
Node Actors Hierarchy



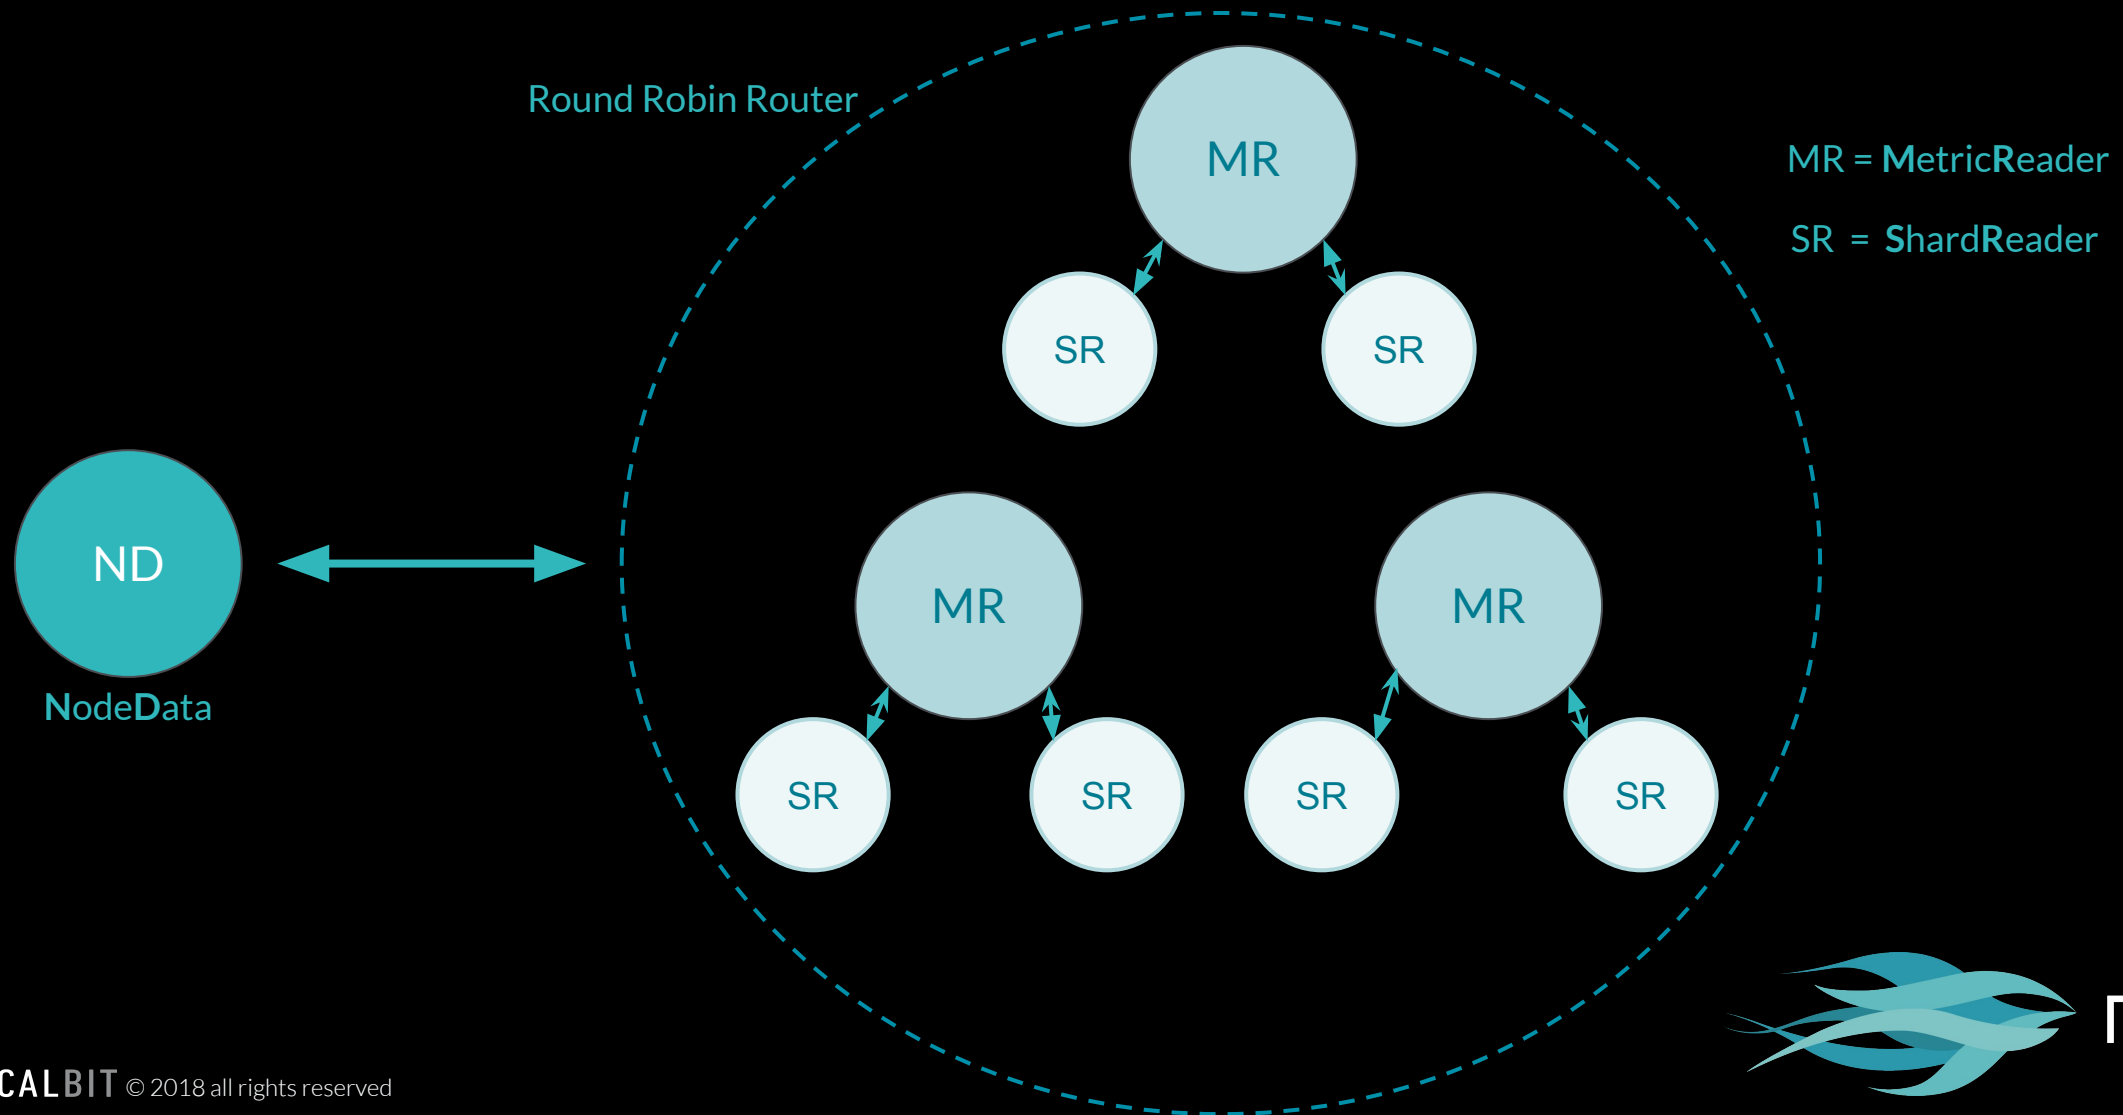
Node Actors Hierarchy - Coordinators



Node Actors Hierarchy - Write Flow



Node Actors Hierarchy - Read Flow (I)



Node Actors Hierarchy - Read Flow (II)

```
private def gatherShardResults(actors: Seq[(Location, ActorRef)], statement: SelectSQLStatement, schema: Schema)(
  postProcFun: Seq[Bit] => Seq[Bit]): Future[Seq[SelectStatementFailed] Either Seq[Bit]] = {

  def sequence[F, A](x: Seq[Either[F, Seq[A]]]): Either[Seq[F], Seq[A]] =
    x partition { _.isLeft } match {
      case (Seq(), r) => Right(r flatMap { _.right.get })
      case (l, _)    => Left(l map { _.left.get })
    }

  Future
    .sequence(actors.map {
      case (_, actor) =>
        (actor ? ExecuteSelectStatement(statement, schema, actors.map(_._1)))
          .mapTo[SelectStatementFailed Either Seq[Bit]]
    })
    .map { rawResults: Seq[Either[SelectStatementFailed, Seq[Bit]]] =>
      sequence(rawResults).map(postProcFun)
    }
}
```



Data Streaming

- Once a new bit is received, it's being sent to *PublisherActor*.
- If the bit matches a registered query it's sent on the corresponding WebSocket via Akka Stream flow.

Problem: unbalance in term of number and frequency between subscription commands and published bits received by *PublisherActor*.

Solution: Akka *UnboundedControlAwareMailbox* implementing a priority queue for command messages.



Akka Cluster Overview

Akka Cluster

Akka Cluster extensions


Akka Distributed Data

Akka Distributed Publish Subscribe



Akka Cluster (I)

“A set of nodes joined together through a membership service”



Akka Cluster (II)

- P2P
- Gossip protocol and failure detection
- Event based notification
- Metrics Collector
- Useful Extensions



Akka Distributed Data

- Akka Distributed Data is useful when you need to share data between nodes in an Akka Cluster.
- It is designed as a key-value store, where the values are Conflict Free Replicated Data Types (CRDTs).
- Supports many data types (Set, Map, Counter etc.)
- Supports different consistency levels for writes and reads
- It's not designed to handle big data



Akka Distributed Publish Subscribe

- Actors can subscribe to a named topic
- Messages are published to a named topic
- The message will be delivered to all subscribers of the topic
- Each node interact with the *DistributedPubSubMediator*
- *At most once* delivery guarantee

```
val mediator = DistributedPubSub(context.system).mediator
mediator ! Subscribe(METADATA_TOPIC, metadataActor)
mediator ! Subscribe(SCHEMA_TOPIC, schemaActor)
mediator ! Subscribe(NODE_GUARDIANS_TOPIC, nodeActorsGuardian)

mediator ! Publish(NODE_GUARDIANS_TOPIC, GetMetricsDataActors)
```



Distributed Design

Overall Architecture

State Replication

Data Replication

Distributed Write Model

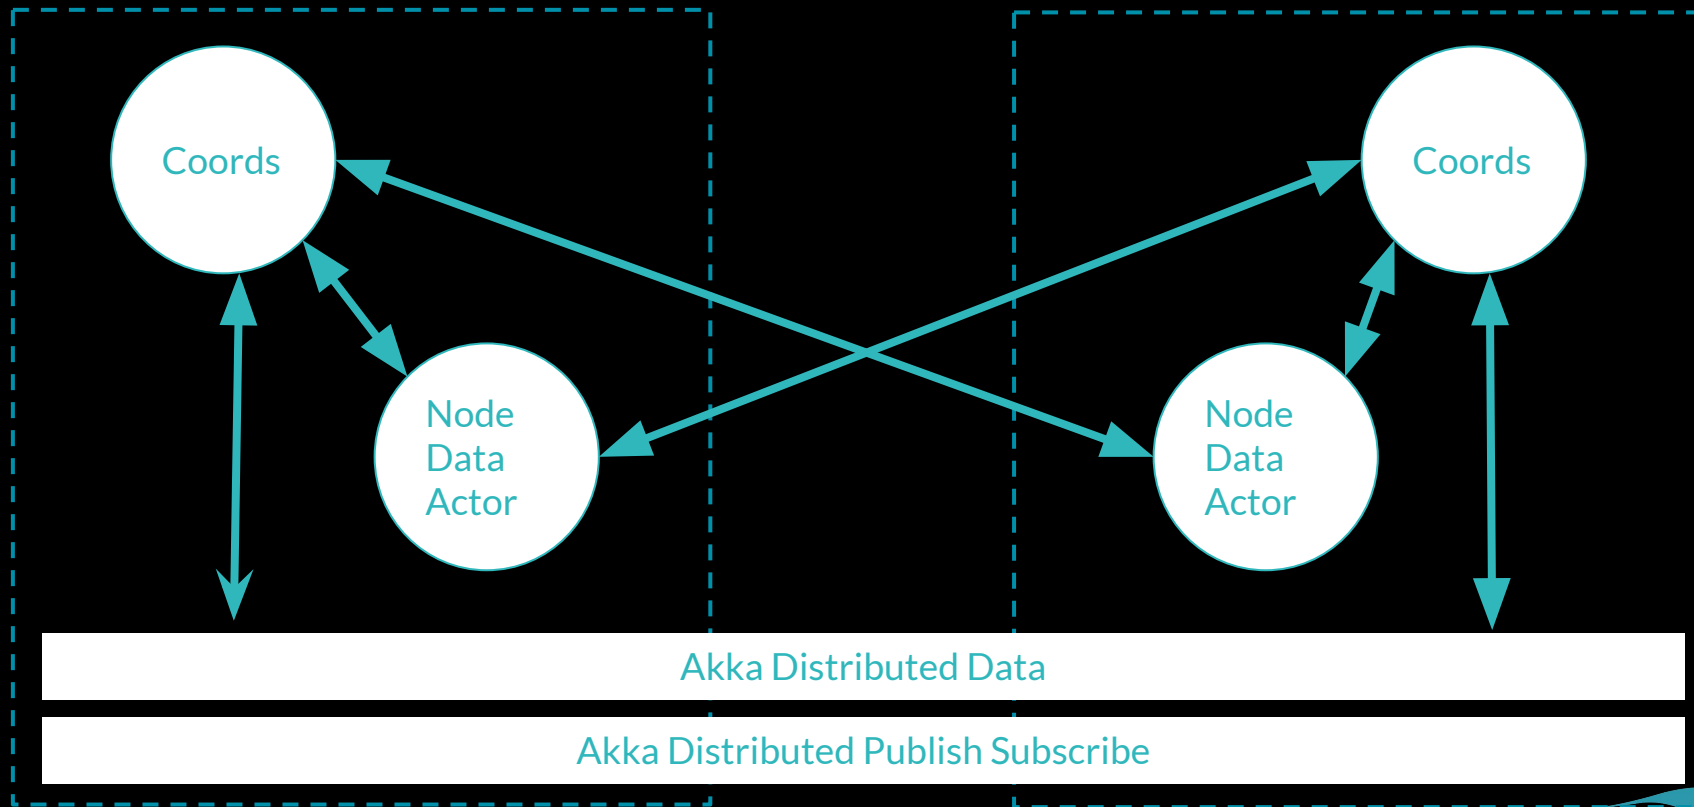
Distributed Read Model

Error Management



Overall Architecture

- **Multimaster replication**, each node can read and write data



Heartbeat protocol

- Leverages Distributed Publish Subscribe
- Every Coordinator is subscribed to a dedicated topic as well as the guardians
- A cluster singleton actor periodically asks guardians to send their data actors reference.
- Cluster events trigger delta updates spread:
 - if a node joins, an add event is disseminated
 - if a node leaves, a remove event is disseminated



State Replication

State = shards locations + schemas



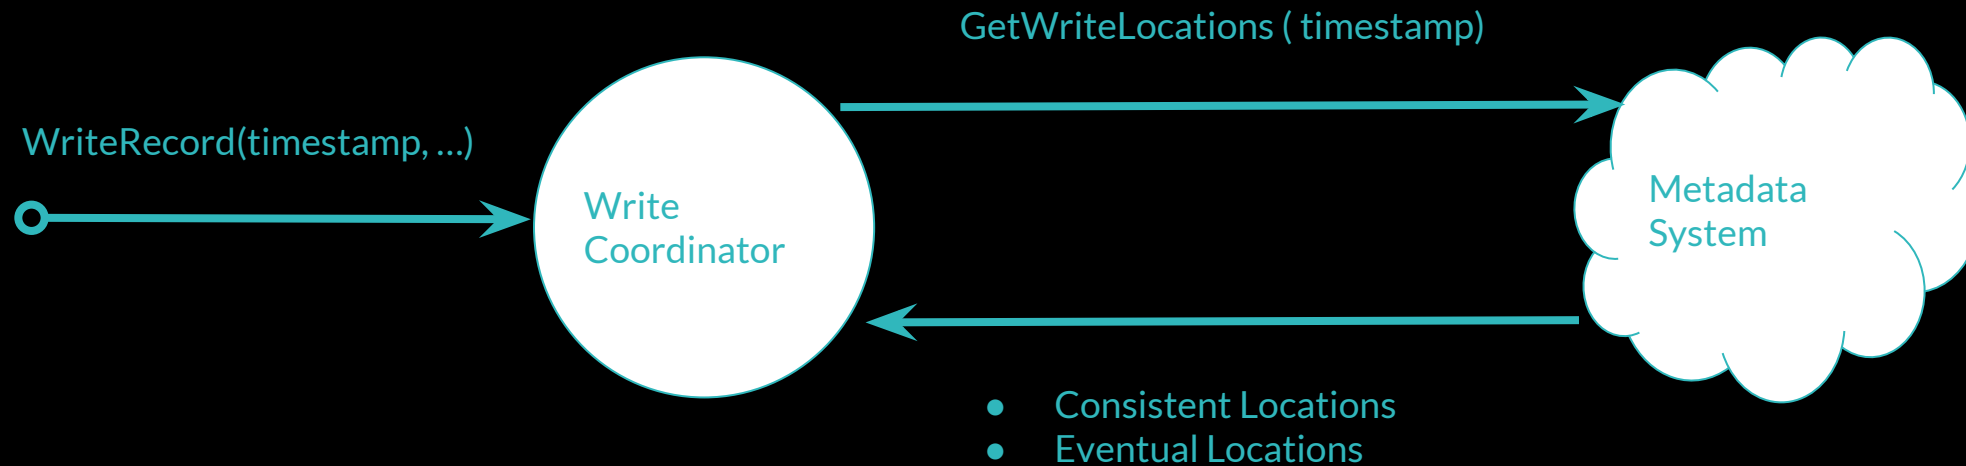
Data Replication

- **Active-active replication** approach
- NSDb implements two levels of replicas in terms of consistency
 - **Consistent replicas:** A record must be correctly acknowledge to all those nodes before the ack can be returned to the caller
 - **Eventual replicas:** the records will be written asynchronously (it fails silently)




Distributed Write Model (I)

1. Record validation
2. Consistent and eventual write locations gathering



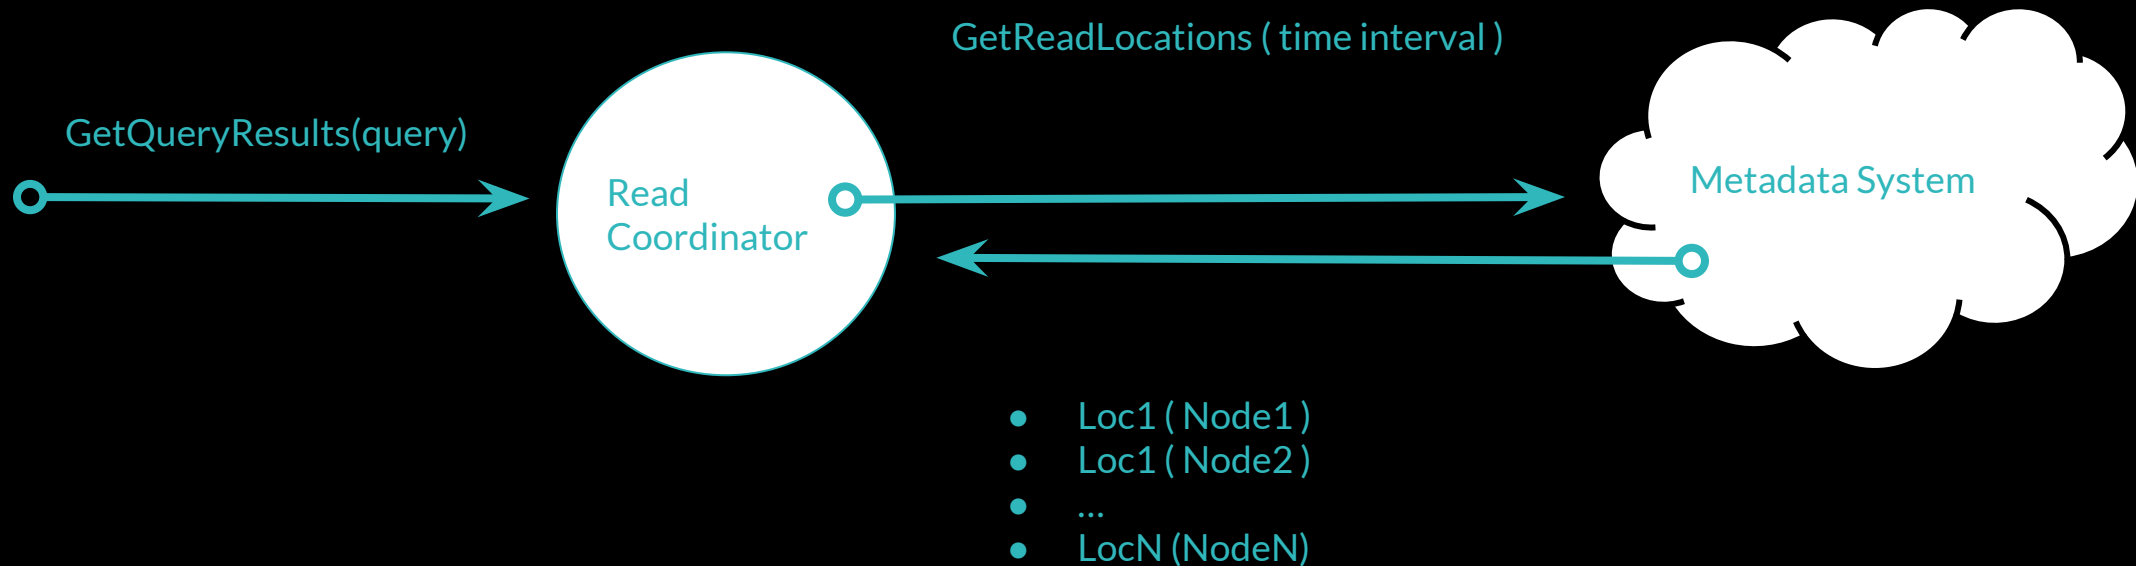
Distributed Write Model (II)

3. Data on Consistent locations written and acknowledge returned to the caller
4. Silently, writes on eventual locations performed



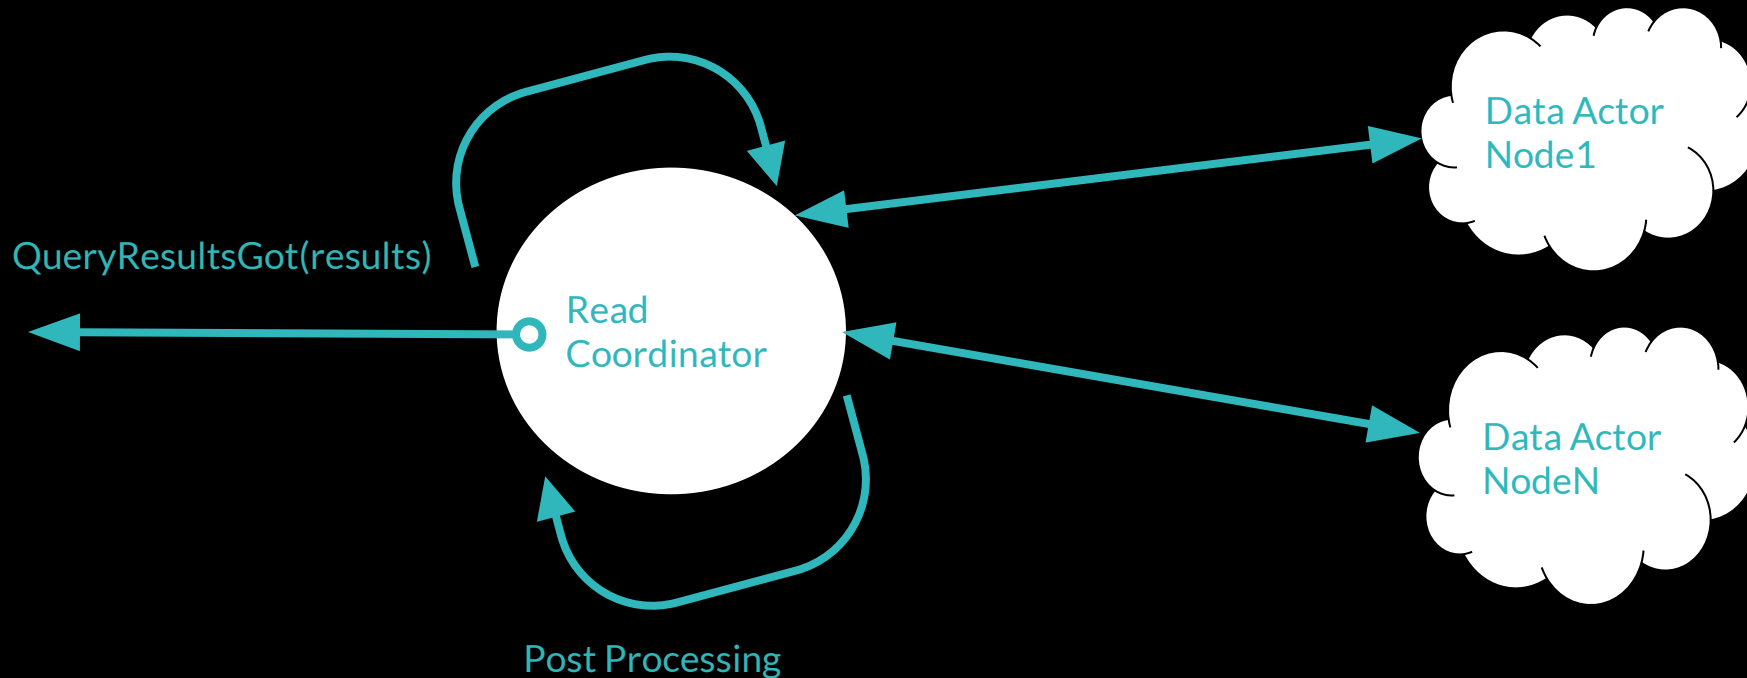
Distributed Read Model (I)

1. Extract time interval from input query where condition (if present)
2. Get locations from metadata system



Distributed Read Model (II)

3. Reduce location lists to one per location
4. Nodes results retrieving (parallel requests to every Node)
5. Post Processing and return result




Error Management (I)

- Write to a set of replicas == distributed transaction
- No isolation
- Saga pattern is applied


Error Management (II)

Two-Phased Commit

Phase 1



Phase 2



credits: @victorklang



Roadmap



Community Edition

NSDb is released under :
Apache 2 License

Reach us on :

<https://github.com/radicalbit/NSDb>



Enterprise Edition

- Support
 - Security
 - OpenID and OAuth support
 - Kerberos Support
 - Metric Versioning
- 

Q&A



GRAZIE!

<radicalbit.team/>
info@radicalbit.io

